

# Algorithm Engineering

False Sharing, Race Conditions, and Schedules

Mark Blacher

Friedrich Schiller University Jena

Winter Semester of 2022/23

# Your Project (Best Case)

- ▶ version control with **Git**
- ▶ the use of **CMake** to build the project
- ▶ correct, efficient, parallel, vectorized, cache-friendly, well organized, **commented**, (portable) C++/C code
- ▶ If you write a **library**, the **API** should be **clear** and **well documented**
- ▶ If you create a **standalone executable**, **document well** the **usage** and the **available command line options** (command line parameters can be parsed with getopt)
- ▶ **unit tests** (for example with **Catch2**)

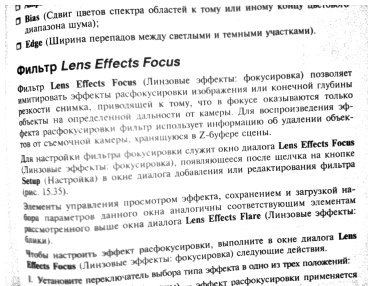

# Project 1: Enhancer for Scanned Images

**Input:** portable pix map (ppm) color image(s) in ASCII-format

**Output:** cleaned ppm image(s) in ASCII-format

**Motivation:** free tool (**executable**) for improving the quality of scanned documents before printing (or sending) them

**Keywords:** noise reduction, background removal, image filter



# Project 2: Generating 2D Datasets With Similar Statistical Properties

**Input:** a 2D point cloud

**Output:** a 2D point cloud with same statistical properties (for a given shape)

**Motivation:** the current approach is slow (written in Python); perhaps it can be done in a more general way than through hordecoded shapes in the program

**Keywords:** simulated annealing, genetic algorithms, mean, standard deviation, correlation

**For more information see:**

<https://www.autodesk.com/research/publications/same-stats-different-graphs>

# Project 3: Parallel In-Place Partitioning, Quickselect and Quicksort 🤖

**Input:** an array of any type of data

**Output:** a correctly modified array (depends on the called function)

**Motivation:** learn advanced parallelization techniques

**Keywords:** partitioning, quickselect, quicksort, atomics (fetch-and-add), header-only library, templates, skewed distributions

**Remark:** Provided Functionality should be similar to `std::partition`, `std::nth_element`, `std::sort` (easy to test the implementation)

**For more information see:**

[A Simple, Fast Parallel Implementation of Quicksort](#)

[How to use templates for sorting](#)

[Parallel implementations for benchmarking](#)

# Your Own Project?

- ▶ You can choose one of three mentioned projects
- ▶ But, it is also possible to come up with your own project
- ▶ If you have an **own project idea, describe it to me in an e-mail**
- ▶ If your own project is suitable for the course, then you can do it

Send me an **e-mail until November 18rd** with a **description of your own project, or, with the information which of the three presented projects** you would like to do.

It might be beneficial for you to **add me to your repository.**

[mark.blacher@uni-jena.de](mailto:mark.blacher@uni-jena.de)

# Avoid False Sharing Among Threads


## Cache Line

The smallest unit of memory that can be transferred between the main memory and the cache (**typically 64 bytes**).

## False Sharing

False sharing occurs when **threads on different processors modify variables that reside on the same cache line**. This invalidates the cache line and forces a memory update to maintain cache coherency.

# False Sharing




[Click here for more information](#)




# Code Example With False Sharing



false\_sharing.cpp 

```
1  #include <iostream>
2  #include <omp.h>
3
4  using namespace std;
5
6  int main() {
7      const int threads = 4; // number of threads
8      volatile long arr[threads]; // no compiler optimizations on arr
9      double start_time = omp_get_wtime(); // wall clock time in seconds
10     #pragma omp parallel num_threads(threads)
11     {
12         const int thread_id = omp_get_thread_num();
13         for (long i = 0; i < 1000000000; ++i) {
14             // each thread uses (probably) the same cache line
15             ++arr[thread_id];
16         }
17     }
18     // cout the byte where arr starts in the cache line
19     cout << ((long) &arr[0]) % 64 << endl;
20     cout << omp_get_wtime() - start_time << " seconds\n";
21 }
```

# Code Example Without False Sharing

no\_false\_sharing.cpp 

```
1  #include <iostream>
2  #include <omp.h>
3
4  using namespace std;
5
6  int main() {
7      const int threads = 4; // number of threads
8      const int PAD = 8; // padding value to avoid false sharing
9      volatile long arr[threads * PAD]; // no compiler optimizations on arr
10     double start_time = omp_get_wtime(); // wall clock time in seconds
11     #pragma omp parallel num_threads(threads)
12     {
13         const int thread_id = omp_get_thread_num();
14         for (long i = 0; i < 1000000000; ++i) {
15             // each thread increments a value on a different cache line
16             ++arr[thread_id * PAD];
17         }
18     }
19     cout << omp_get_wtime() - start_time << " seconds\n";
20 }
```

# Prevent Race Conditions With Mutual Exclusion

## Race Condition

A race condition occurs when **two or more threads try to change shared data** at the same time.

## Mutual Exclusion

A form of synchronization in which **only one thread** at a time can execute a block of code.




# Mutual Exclusion Constructs

- ▶ `#pragma omp critical` *// critical block runs on one thread at a time*
- ▶ `std::mutex` *// an instance of std::mutex should be shared between threads*
- ▶ `#pragma omp atomic` *// updates a memory location atomically*
- ▶ ...

Usually a **thread** that reaches a mutual exclusion construct **waits until it gains exclusive access**.

# Can You Spot the Race Condition Bug?



race\_condition.cpp 

```
1  double pi = 0.0;
2  #pragma omp parallel num_threads(4)
3  {
4      int num_threads = omp_get_num_threads();
5      int thread_id = omp_get_thread_num();
6      double sum_local = 0.0; // for summing up heights locally
7      for (int i = thread_id; i < num_steps; i += num_threads) {
8          double x = (i + 0.5) * width; // midpoint
9          sum_local = sum_local + (1.0 / (1.0 + x * x)); // add new height
10     }
11     pi += sum_local * 4 * width;
12 }
```

# Why Is This Code Slow?



pi\_slow.cpp 

```
1  double pi = 0.0;
2  #pragma omp parallel num_threads(4)
3  {
4      int num_threads = omp_get_num_threads();
5      int thread_id = omp_get_thread_num();
6      for (int i = thread_id; i < num_steps; i += num_threads) {
7          double x = (i + 0.5) * width; // midpoint
8          #pragma omp atomic
9              pi = pi + (4.0 / (1.0 + x * x)) * width;
10     }
11 }
```

# The `#pragma omp for` Construct

`#pragma omp for` splits up a for loop within a parallel region among the threads in a team

`pragma_omp_for.cpp` 

```
1  image *denoise_image(image *in) { // mock denoise
2      image *out = nullptr;
3      for (volatile int i = 0; i < 20000000; ++i) {}
4      return out;
5  }
6
7  int main() {
8      int amount_images = 64;
9      vector<image *> images(amount_images, nullptr); // mock images
10     vector<image *> denoised_images(amount_images); // mock output
11     double start = omp_get_wtime();
12     #pragma omp parallel
13     {
14         #pragma omp for // denoise a bunch of images
15         for (int i = 0; i < amount_images; ++i) {
16             denoised_images[i] = denoise_image(images[i]);
17         }
18     }
19     cout << omp_get_wtime() - start << " seconds" << endl;
20 }
```

# Unbalanced Workload

unbalanced\_workload.cpp 

```
1 image *denoise_image(image *in, int scale) { // mock denoise
2     image *out = nullptr;
3     const int loop_until = 20000 * scale; // unbalanced workload
4     for (volatile int i = 0; i < loop_until; ++i) {}
5     return out;
6 }
7
8 int main() {
9     int amount_images = 64;
10    vector<image *> images(amount_images, nullptr); // mock images
11    vector<image *> denoised_images(amount_images); // mock output
12    double start = omp_get_wtime();
13
14    // #pragma omp parallel for is an "OpenMP shortcut"
15    #pragma omp parallel for num_threads(4) // denoise a bunch of images
16    for (int i = amount_images - 1; i > -1; --i) {
17        denoised_images[i] = denoise_image(images[i], i * i);
18    }
19
20    cout << omp_get_wtime() - start << " seconds" << endl;
21 }
```



# The Schedule Clause

The schedule clause affects **how loop iterations are mapped onto threads when using `#pragma omp for`**

By far the **most important schedule clauses** are:

- ▶ **static** (we make the scheduling decision at compile time)
- ▶ **dynamic** (deciding at runtime, how to break up work between the threads)

## Usage:

- ▶ `schedule(static [,chunk])`  
Deal-out blocks of iterations of size "chunk" to each thread.
- ▶ `schedule(dynamic [,chunk])`  
Each thread grabs "chunk" iterations off a queue until all iterations have been handled.

# Schedule Clause Demo

schedule\_clause.cpp 

```
1 int main() {
2     #pragma omp parallel num_threads(4)
3     {
4         // each thread has its own stringstream instance
5         stringstream thread_iterations;
6         // task for you: experiment with different schedules
7         #pragma omp for schedule(static, 8)
8         for (int i = 0; i < 32; ++i) {
9             // write the iterations of a thread to the stream
10            thread_iterations << i << " ";
11        }
12        // only one thread is allowed to output to the console
13        #pragma omp critical
14            cout << "thread " << omp_get_thread_num()
15                << " computed the iterations: "
16                << thread_iterations.str() << endl;
17    }
18 }
```

# Choosing the Right Schedule Clause

If each iteration in the for loop needs the **same workload**, make the **schedule static**.

If the **workload** is **unbalanced** between iterations, make the **schedule dynamic**.

# How to Stop a Parallel For Loop

- ▶ Unfortunately, **you can't *easily* just break out** of a **parallel for loop** in OpenMP
- ▶ But there's a **workaround**

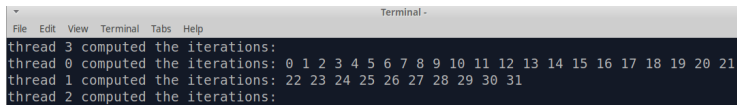
# Don't Stop a Parallel For Loop, Just Avoid Work

continue\_hack.cpp 

```
1  bool is_solution(int number) { // test if number solves the problem
2      for (volatile int i = 10000000; i--;) {} // mock computation
3      return number > 133 && number < 140;
4  }
5
6  int main() {
7      constexpr int biggest_possible_number = 10000;
8      // To avoid undefined code, it is necessary to use std::atomic for variables
9      // where race conditions may happen.
10     // https://databasearchitects.blogspot.com/2020/10/c-concurrency-model-on-x86-for-dummies.html
11     atomic<bool> solution_found(false); // if true then we found the solution
12     atomic<int> final_solution(INT32_MAX);
13     const double start = omp_get_wtime();
14
15     #pragma omp parallel for schedule(dynamic) // start parallel region
16     for (int i = 0; i < biggest_possible_number; ++i) {
17         if (solution_found) // we found the solution, just continue iterating
18             continue;
19         if (is_solution(i)) { // find some solution, not necessary the smallest
20             solution_found = true;
21             final_solution = i;
22         }
23     } // end parallel region
24     // check if we've found a solution at all is omitted, you can add the check
25     cout << "The solution is: " << final_solution << endl;
26     cout << omp_get_wtime() - start << " seconds" << endl;
27 }
```

# Coding Warmup

1. Fix the **race condition bug** on page 13 with a **`std::mutex`**.
2. **Reduce** the **runtime** of the image denoising program on slide 16 by adding an **appropriate schedule**.
3. What **schedule** on slide 18 produces the **following (bad) pattern?**

A terminal window titled "Terminal -" with a menu bar (File, Edit, View, Terminal, Tabs, Help). The output shows four lines of text: "thread 3 computed the iterations:", "thread 0 computed the iterations: 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21", "thread 1 computed the iterations: 22 23 24 25 26 27 28 29 30 31", and "thread 2 computed the iterations:".

```
thread 3 computed the iterations:
thread 0 computed the iterations: 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21
thread 1 computed the iterations: 22 23 24 25 26 27 28 29 30 31
thread 2 computed the iterations:
```

4. **Modify** the **program** on slide 21 so that it **always** finds the **smallest possible solution**.

# Solution for the Warmup *Smallest* Program (Lock)

continue\_hack\_min\_solution\_lock.cpp 

```
1  bool is_solution(int number) { // test if number solves the problem
2      for (volatile int i = 10000000; i--;) {} // mock computation
3      return number > 133 && number < 140;
4  }
5
6  int main() {
7      constexpr int biggest_possible_number = 10000;
8      atomic<int> final_solution(INT32_MAX);
9      const double start = omp_get_wtime();
10     mutex m; // mutual exclusion construct for updating final_solution
11
12     #pragma omp parallel for schedule(dynamic) // we could also use schedule(static, 1)
13     for (int i = 0; i < biggest_possible_number; ++i) {
14         if (final_solution < i) // we found a smaller solution, just continue iterating
15             continue;
16         if (is_solution(i)) { // one solution found, update final_solution if it's smaller
17             lock_guard<mutex> lg(m); // constructs a lock_guard, locking the given mutex
18             final_solution = min(i, final_solution.load()); // only one thread computes min
19         }
20     }
21     // end parallel region
22     // check if we've found a solution at all is omitted, you can add the check
23     cout << "The smallest solution is: " << final_solution << endl;
24     cout << omp_get_wtime() - start << " seconds" << endl;
25 }
```

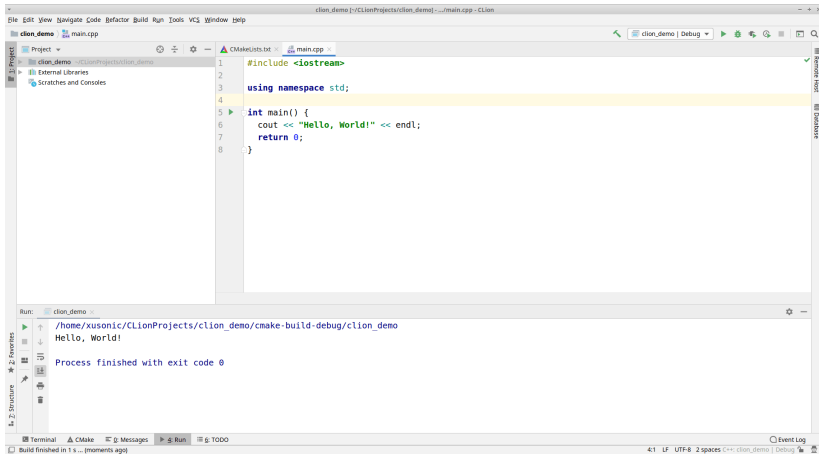
# Solution for the Warmup *Smallest* Program (Atomic)

continue\_hack\_min\_solution\_atomic.cpp 

```
1  bool is_solution(int number) { // test if number solves the problem
2      for (volatile int i = 10000000; i--;) {} // mock computation
3      return number > 133 && number < 140;
4  }
5
6  int main() {
7      constexpr int biggest_possible_number = 10000;
8      atomic<int> final_solution(INT32_MAX); // now we use an atomic
9      const double start = omp_get_wtime();
10
11     #pragma omp parallel for schedule(dynamic) // start parallel region
12     for (int i = 0; i < biggest_possible_number; ++i) {
13         if (final_solution < i) // we found a smaller solution, just continue iterating
14             continue;
15         if (is_solution(i)) { // one solution found, update final_solution if it's smaller
16             int previous = final_solution.load();
17             while (previous > i && !final_solution.compare_exchange_weak(previous, i)) {}
18             // http://www.cplusplus.com/reference/atomic/atomic/compare_exchange_weak/
19             // atomic.compare_exchange_weak(T& expected, T desired);
20             // 1. Compare a given value expected with the value stored in atomic.
21             // 2. If both values stored in expected and atomic coincide then set atomic to a
22             // given value desired, otherwise write the actual value stored in atomic to expected.
23             // 3. Return true if the swap in 2. was successful, otherwise return false.
24         }
25     } // end parallel region
26     // check if we've found a solution at all is omitted, you can add the check
27     cout << "The solution is: " << final_solution << endl;
28     cout << omp_get_wtime() - start << " seconds" << endl;
29 }
```



# Clion Demo



# Exam Assignments

- ▶ What causes **false sharing**?
- ▶ How do **mutual exclusion constructs** prevent **race conditions**?
- ▶ Explain the differences between **static** and **dynamic schedules in OpenMP**.
- ▶ What can we do if we've **found a solution** while running a **parallel for loop** in OpenMP, but still have **many iterations left**?
- ▶ Do the **coding warmup** on **slide 22**. Explain in your own words how `std::atomic::compare_exchange_weak` work.
- ▶ **Rewrite the parallel program** for estimating  $\pi$  from the last exercise  $\Rightarrow$  to use the construct `#pragma omp for`.