

Algorithm Engineering

Ordered Clause, Collapse Clause, Reductions, Barriers, and
Storage Attributes


Mark Blacher

Friedrich Schiller University Jena

Winter Semester of 2022/23

The ordered Clause

Different threads execute concurrently until they encounter the **ordered region**, which is then **executed sequentially** in the **same order** as it would get executed in a **serial loop**.

ordered.cpp 

```
1  int main() { // find smallest solution with the ordered clause
2      constexpr int biggest_possible_number = 10000;
3      atomic<bool> solution_found(false); // if true then we found the solution
4      int final_solution = INT32_MAX;
5      const double start = omp_get_wtime();
6      #pragma omp parallel for ordered schedule(dynamic) // start parallel region
7      for (int i = 0; i < biggest_possible_number; ++i) {
8          if (solution_found) // we found the solution, just continue iterating
9              continue;
10         if (is_solution(i)) {
11             #pragma omp ordered // ordered region
12             if (!solution_found) { // ordered execution of if statement
13                 solution_found = true; // no race condition
14                 final_solution = i;
15             }
16         }
17     } // end parallel region
18     cout << "The solution is: " << final_solution << endl;
19     cout << omp_get_wtime() - start << " seconds" << endl;
20 }
```

Cancellation Points

cancellation_points.cpp 

```
1 // actually, we can stop a parallel for loop with cancellation points, but it's a little complicated
2 int main(int argc, char **argv) { // since OpenMP 4.0 we have cancellation points
3     if (!omp_get_cancellation()) { // if no cancellations enabled, enable it
4         cout << "Enabling cancellation and rerunning program\n" << endl;
5         const char* enable_cancellation = "OMP_CANCELLATION=true";
6         // const_cast can be used to pass const data to a function that doesn't receive const
7         putenv(const_cast<char *>(enable_cancellation)); // set cancellation environment variable
8         execv(argv[0], argv); // rerun program, because can't enable cancellations in the program itself
9     } // execv replaces the current process image with a new process image
10    constexpr int biggest_possible_number = 10000;
11    atomic<int> final_solution(INT32_MAX);
12    const double start = omp_get_wtime();
13
14    #pragma omp parallel // start parallel region
15    {
16        #pragma omp for schedule(dynamic)
17        for (int i = 0; i < biggest_possible_number; ++i) {
18            if (is_solution(i)) { // find some solution, not necessary the smallest
19                final_solution = i;
20            }
21        }
22        #pragma omp cancellation point for // signal cancellation, because we found a solution
23        // cancel for loop if cancellations signalled
24    } // end parallel region
25
26    // check if we've found a solution at all is omitted, you can add the check
27    cout << "The solution is: " << final_solution << endl;
28    cout << omp_get_wtime() - start << " seconds" << endl;
29 }
30 }
```


Nested For Loops

We can **parallelize nested for loops** with the **collapse** clause.

collapse.cpp 

```
1  int main() {
2      const int N = 10;
3      const int M = 12;
4
5      /**
6       * will form a single loop of length  $N * M = 120$ 
7       * and then parallelize that
8       *
9       * can be useful for balancing the work
10     */
11     #pragma omp parallel for collapse(2)
12     for (int i = 0; i < N; i++) {
13         for (int j = 0; j < M; j++) {
14             // do useful work with i and j
15         }
16     }
17 }
```

Collapse Clause Example

grid_search.cpp 

```
1 // mock k-nearest neighbors classification
2 void knn(int k, const string &weight, const string &metric) {
3     // larger k needs more runtime
4     for (volatile int i = 0; i < 10000000 * k; ++i) {}
5     // computing accuracy of classification is omitted here
6     #pragma omp critical // output which thread did what
7     cout << "k: " << k << ", weight: " << weight << ", metric: " << metric
8         << ", computed with thread: " << omp_get_thread_num() << endl;
9 }
10
11 int main() {
12     vector<int> ks{1, 3, 5, 7, 9, 11};
13     vector<string> weights{"uniform", "distance"};
14     vector<string> metrics{"euclidean", "manhattan"};
15     const double start = omp_get_wtime();
16     // parallel grid search for tuning the hyperparameters in knn
17     #pragma omp parallel for collapse(3) schedule(dynamic)
18     for (uint64_t i = 0; i < ks.size(); ++i)
19         for (uint64_t j = 0; j < weights.size(); ++j)
20             for (uint64_t k = 0; k < metrics.size(); ++k)
21                 knn(ks[i], weights[j], metrics[k]);
22
23     cout << omp_get_wtime() - start << " seconds" << endl;
24 }
```

Reduction

Reduction is an **associative** and **commutative operation**. It is used in parallel programming to **reduce many values** into a **single result**.

In OpenMP the reduction clause looks like this:

`reduction(op:list)`

`op` can be `+`, `*`, `-`, `min`, `max`, `&`, `|`, `^`, `&&` and `||`

`list` contains the variables, separated by commas, that are to be reduced

Reduction Example

reduction.cpp 

```
1 int main() {
2     int n = 100000000;
3     // create a vector of size n, all values 1.0
4     // long double is "usually" a 128 bit float data type
5     // L indicates that 1.0 is a long double literal
6     vector<long double> vec(n, 1.0L);
7     long double sum = 0.0L;
8
9     const double start = omp_get_wtime();
10
11    #pragma omp parallel for reduction(+: sum)
12    for (int i = 0; i < n; ++i) {
13        sum += vec[i]; // sum up all values of vec
14    }
15
16    cout << fixed << "sum: " << sum << endl;
17    cout << omp_get_wtime() - start << " seconds" << endl;
18 }
```


How Reductions Work Internally

- ▶ A local copy of each *list* variable is made and initialized depending on the *op* (0 for +)
- ▶ Updates occur on the local copy
- ▶ Local copies are reduced into a single value and combined with the original global value

Synchronization With Barriers

Barrier

A barrier means that **any thread must stop at this point and cannot proceed until all other threads reach this barrier.**

barrier.cpp 

```
1 int main() {
2 #pragma omp parallel
3 {
4     stringstream info;
5     info << "Hello from thread " << omp_get_thread_num() << endl;
6     cout << info.str();
7 #pragma omp barrier // move on as soon as all threads printed Hello ...
8     info.str(""); // "clear" stringstream variable
9     info << "Goodbye from thread " << omp_get_thread_num() << endl;
10    cout << info.str();
11 }
12 }
```

Implicit Barriers in OpenMP


The **for**, **sections**, **single** and **parallel** constructs in OpenMP have an **implicit barrier**.

Implicit barrier means that **there is a barrier without you explicitly positioning it there.**

In the constructs **for**, **sections** and **single** you can remove the implicit barrier with the **nowait** clause.

You **can't remove the implicit barrier** in the **parallel** construct.

Sections

sections.cpp 

```
1  #include <iostream>
2
3  using namespace std;
4
5  int main() {
6  #pragma omp parallel
7      {
8  #pragma omp sections // add nowait to remove the implicit barrier
9      { // Each section is executed once by one of the threads in the team
10 #pragma omp section // some thread executes this block of code
11     { cout << "Hello from section 1\n"; }
12 #pragma omp section // some thread executes this block of code
13     { cout << "Hello from section 2\n"; }
14 #pragma omp section // some thread executes this block of code
15     { cout << "Hello from section 3\n"; }
16     } // after the sections construct is an implicit barrier
17     cout << "Hello after the sections\n"; // this line is executed by
18                                           // every thread in the team
19 }
20 }
```

Single

The **single construct** specifies that the **associated structured block** is **executed by only one of the threads** in the team (not necessarily the master thread).

single.cpp 

```
1  #include <iostream>
2  #include <omp.h>
3
4  using namespace std;
5
6  int main() {
7      #pragma omp parallel
8      {
9          #pragma omp single // only one thread executes the code in the single
10         {
11             cout << "from single: " << omp_get_thread_num() << "\n";
12         } // implicit barrier after single
13         cout << "after single\n";
14     };
15 }
```

Useful Runtime Library Routines


These are the **routines** you already know:

- ▶ `omp_set_num_threads(int)` *// setting desired number of threads in the parallel region*
- ▶ `omp_get_num_threads()` *// number of threads*
- ▶ `omp_get_thread_num()` *// thread id*
- ▶ `omp_get_wtime()` *// wall-clock time in seconds*

Here are **some more routines**:

- ▶ `omp_get_num_procs()` *// number of logical cores*
- ▶ `omp_get_max_threads()` *// maximum number of threads in a parallel region*
- ▶ `omp_in_parallel()` *// true if called inside a parallel region*

Library Routines Demo

library_routines.cpp 

```
1 int main() {
2     cout << "num_threads: " << omp_get_num_threads() << "\n"; // prints 1
3     // but, if we call omp_get_num_threads in a parallel region
4     // then we get the number of threads that are executed inside the parallel region
5     #pragma omp parallel
6     #pragma omp single
7     cout << "num_threads: " << omp_get_num_threads() << "\n\n"; // e.g 4
8
9     // maximum number of threads in a parallel region (e.g. 4)
10    cout << "max_threads: " << omp_get_max_threads() << "\n"; // prints 4
11    omp_set_num_threads(2); // change amount of threads in a parallel region
12    cout << "max_threads: " << omp_get_max_threads() << "\n\n"; // now it's 2
13
14    // prints 4, if my machine has 4 logical cores
15    cout << "num_procs: " << omp_get_num_procs() << "\n\n";
16    // logical cores are the number of physical cores times the number
17    // of threads (called "hyperthreads") that can run on each core
18
19    cout << "in_parallel: " << omp_in_parallel() << "\n"; // prints 0 ("false")
20    // if we call omp_in_parallel in a parallel region then we get 1 ("true")
21    #pragma omp parallel
22    #pragma omp single
23    cout << "in_parallel: " << omp_in_parallel() << "\n";
24 }
```

Storage Attributes

Storage attributes determine how **variables defined outside the parallel region** are **handled** ("stored") **inside the parallel region**. *// simplified explanation but enough for our purposes*

Useful storage attributes:

- ▶ **shared** *// this is default, the variable is shared between threads*
- ▶ **private** *// create uninitialized copy of the variable for each thread*
- ▶ **firstprivate** *// create initialized one-to-one copy of the variable for each thread*

Storage Attributes Demo

private_vs_firstprivate.cpp 

```
1 int main() {
2     int x = 10;
3     omp_set_num_threads(4);
4
5     cout << "private(x): ";
6     #pragma omp parallel private(x)
7     { // x is an uninitialized local copy in each thread
8         x += omp_get_thread_num(); // (probably) 0 + thread_num
9     }
10    #pragma omp critical
11        cout << x << " "; // (probably) something like: 3 2 1 0
12
13    cout << "\nfirstprivate(x): ";
14    #pragma omp parallel firstprivate(x)
15    { // x is an initialized local one-to-one copy in each thread
16        x += omp_get_thread_num(); // 10 + thread_num
17    }
18    #pragma omp critical
19        cout << x << " "; // something like: 10 12 11 13
20
21    cout << "\n'global' x: " << x << endl; // x = 10 (global x didn't change)
22 }
```


Storage Attributes Remarks

- ▶ Pragas like this one:

```
#pragma omp parallel default(none) shared(a, b) firstprivate(x, my_vector)
```

are common

- ▶ **default(none)** is good for debugging, since you must consciously specify the storage attribute for each 'global' variable used within the parallel region
- ▶ Be careful if you use **private(my_vector)**, each thread inside the parallel region creates a vector object of size=0

segmentation_fault.cpp ➡

```
1 int main() {  
2     vector<int> my_vector{1, 2, 3};  
3     #pragma omp parallel default(none) shared(cout) private(my_vector)  
4     {  
5         // program crashes since my_vector is uninitialized (size = 0)  
6         cout << my_vector[0]; // segmentation fault  
7     }  
8 }
```



Coding Warmup

1. **Parallelize** the **serial π program** \Rightarrow from the first lecture by **adding** only **one extra line** of code.
2. Write a **parallel program** for computing π using **only standard C++** (no OpenMP).
3. The program on page 20 generates and saves an image of the mandelbrot set. **Reduce** the **runtime** of the program **without degrading** the **image quality**.

Compile the programs with at least the following **flags**:

-Ofast -std=c++11 -march=native -fopenmp

OpenMP vs C++ Threads

pi_openmp.cpp ➡

```
1 #pragma omp parallel for reduction(+:sum)
2   for (int i = 0; i < num_steps; i++) {
3       double x = (i + 0.5) * width; // midpoint
4       sum = sum + (1.0 / (1.0 + x * x)); // add new height of a rectangle
5   }
```

pi_cpp_threads.cpp ➡

```
1 // thunk is called from different threads
2 void thunk(int num_threads, int thread_id, int num_steps, double width,
3            double &sum, mutex &m) {
4     double sum_local = 0.0; // for summing up heights locally
5     for (int i = thread_id; i < num_steps; i += num_threads) {
6         double x = (i + 0.5) * width; // midpoint
7         sum_local = sum_local + (1.0 / (1.0 + x * x)); // add new height
8     }
9     lock_guard<mutex> lg(m); // avoids race condition on sum
10    sum += sum_local; // update global sum
11 }
```

It's good to know how you could parallelize without OpenMP.

Generating an Image

mandelbrot.cpp 

```
1  int main() { // generate mandelbrot pgm (portable graymap)
2      const string image_name = "mandelbrot.pgm";
3      remove(image_name.c_str()); // remove file from disk
4      const double start = omp_get_wtime();
5      ofstream image(image_name); // file output stream
6      if (image.is_open()) {
7          image << "P2\n" << width << " " << height << " 255\n"; // pgm header
8          for (int i = 0; i < height; i++) {
9              for (int j = 0; j < width; j++) {
10                 image << compute_pixel(j, i) << "\n"; // write pixel value
11             }
12         }
13         image.close(); // close file output stream
14     } else {
15         cout << "Could not open the file!";
16     }
17     cout << omp_get_wtime() - start << " seconds" << endl;
18 }
```

Possible Solution for the Warmup Image Exercise

mandelbrot_tuned.cpp ➡

```
1  int main() { // generate mandelbrot pgm (portable graymap)
2      const string image_name = "mandelbrot.pgm";
3      remove(image_name.c_str()); // remove file from disk
4      const double start = omp_get_wtime();
5      vector<string> look_up{256}; // look-up table for strings
6      for (int i = 0; i < 256; ++i) { look_up[i] = to_string(i) + "\n"; } // fill look-up table
7      ofstream image(image_name); // file output stream
8      if (image.is_open()) {
9          image << "P2\n" << width << " " << height << " 255\n"; // write pgm header
10         #pragma omp parallel // start parallel region
11         {
12             string buffer; // each thread has it's own string buffer
13             buffer.reserve(width * 4); // reserve enough space to avoid reallocation
14             #pragma omp for schedule(dynamic) ordered // we could also use schedule(static, 1)
15             for (int i = 0; i < height; i++) {
16                 buffer.clear(); // clear content of buffer
17                 for (int j = 0; j < width; j++) {
18                     buffer += look_up[compute_pixel(j, i)]; // fill buffer
19                 }
20             #pragma omp ordered // write rows of image in sequential order
21             image << buffer; // write string buffer to file
22         }
23     } // end parallel region
24     image.close(); // close file output stream
25 } else {
26     cout << "Could not open the file!";
27 }
28 cout << omp_get_wtime() - start << " seconds" << endl;
29 }
```

Exam Assignments

- ▶ How does the **ordered clause** in OpenMP work in conjunction with a parallel for loop?
- ▶ What is the **collapse clause** in OpenMP good for?
- ▶ Explain how **reductions** work internally in OpenMP.
- ▶ What is the purpose of a **barrier** in parallel computing?
- ▶ Explain the differences between the library routines `omp_get_num_threads()`, `omp_get_num_procs()` and `omp_get_max_threads()`.
- ▶ Clarify how the storage attributes **private** and **firstprivate** differ from each other.
- ▶ Do the **coding warmup** on **slide 18**. Write in **pseudo code** how the **computation of π** can be **parallelized** with **simple threads**.