

Algorithm Engineering

Tasks, Merge Sort, Multithreaded Merging

Mark Blacher

Friedrich Schiller University Jena

Winter Semester of 2022/23

Major OpenMP Constructs We've Covered So Far

- ▶ To **create a team of threads**

```
#pragma omp parallel
```

- ▶ To **share work between threads**

```
#pragma omp for
```

```
    schedule(dynamic), collapse(3), ordered, reduction(+:sum)
```

```
#pragma omp single
```

- ▶ To **prevent conflicts** (prevent races)

```
#pragma omp critical
```

```
#pragma omp atomic
```

```
#pragma omp barrier
```

- ▶ **Storage Attributes**

```
private(a)
```

```
firstprivate(b)
```

```
shared(c)
```

OpenMP Tasks


- ▶ Tasks are **independent units of work**
- ▶ Tasks have **potential to parallelize recursive function calls**
- ▶ A **queuing system dynamically handles the assignment of threads to the tasks** () \Rightarrow threads pick up a task from the queue until the queue is empty

Task Example

task_example.cpp 

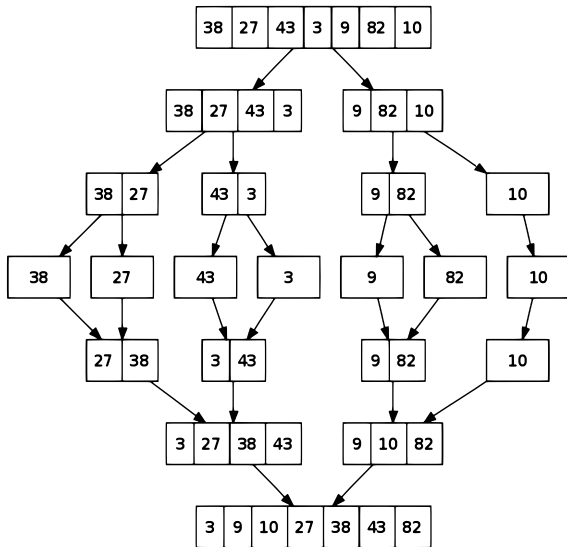
```
1  // There are two tasks. One prints the word "race ".
2  // The other one prints "car ". The order in which tasks are
3  // executed is runtime dependent and may vary across multiple runs.
4  int main() {
5  #pragma omp parallel
6      { // we need the single construct to guarantee that
7        // each task is generated only once
8  #pragma omp single // one thread generates the tasks
9      {
10         // Each #pragma omp task directive defines a task
11  #pragma omp task // define task and put it in the task queue
12         { cout << "race "; } // task 1
13  #pragma omp task // define task and put it in the task queue
14         { cout << "car "; } // task 2
15     } // end of single region (implicit barrier)
16 } // end of parallel region
17 }
```

Taskwait Construct

taskwait.cpp 

```
1  int main() {
2  #pragma omp parallel
3      { // we need the single construct to guarantee that
4        // each task is generated only once
5  #pragma omp single // one thread generates the tasks
6      {
7          cout << "A "; // printed before tasks are generated
8          // Each #pragma omp task directive defines a task
9  #pragma omp task // define task and put it in the task queue
10         { cout << "race "; } // task 1
11  #pragma omp task // define task and put it in the task queue
12         { cout << "car "; } // task 2
13  #pragma omp taskwait // wait here until the tasks are completed
14         cout << "is fun to watch.\n";
15     } // end of single region
16 } // end of parallel region
17 }
```

Merge Sort



Divide and Conquer Example (Merge Sort)

Try to **reduce** the **runtime** of this naive merge sort implementation.

slow_merge_sort.cpp ➡

```
1 void merge_sort_naive(int *arr, int n) { // slow merge sort
2     if (n > 1) { // TODO: use insertion sort for small n
3         const int size_a = n / 2;
4         const int size_b = n - size_a;
5         // TODO: make next recursive call a task
6         merge_sort_naive(arr, size_a); // recursive call
7         merge_sort_naive(arr + size_a, size_b); // recursive call
8         // TODO: here should be a taskwait
9         int *c = new int[n]; // TODO: avoid using heap for small n
10        merge(arr, arr + size_a, c, size_a, size_b, n);
11        memcpy(arr, c, sizeof(int) * n);
12        delete[](c);
13    }
14 }
```

Slightly Improved Merge Sort

merge_sort.cpp 

```
1 void merge_sort(int *arr, int n) {
2     if (n > 1) {
3         if (n < 32) { // insertion sort for n smaller than 32
4             insertion_sort(arr, n);
5             return;
6         }
7         const int size_a = n / 2;
8         const int size_b = n - size_a;
9         #pragma omp task if (size_a > 10000) // make task if size_a > 10000, one task is enough
10        merge_sort(arr, size_a); // parallel recursive call
11        merge_sort(arr + size_a, size_b); // recursive call
12        #pragma omp taskwait // wait until both subarrays are sorted
13        if (n < 8192) { // allocate array on the stack for small n
14            // sometimes called _alloca() or _malloca(), or "int c[n]"
15            int* c = (int*) alloca(n * sizeof(int));
16            merge(arr, arr + size_a, c, size_a, size_b, n);
17            memcpy(arr, c, sizeof(int) * n); // copying can be tuned away...
18            return;
19        }
20        int *c = new int[n];
21        merge(arr, arr + size_a, c, size_a, size_b, n);
22        memcpy(arr, c, sizeof(int) * n); // copying can be tuned away...
23        delete[](c);
24    }
25 }
26
27 void merge_sort_run(int *arr, int n) {
28     #pragma omp parallel
29     #pragma omp single nowait
30     merge_sort(arr, n);
31 }
```


Avoid Copying Data in Merge Sort

merge_sort_no_copying.cpp 

```
1 void merge_sort_helper(int *arr, int *buffer, int n) { // helps to implement no copying strategy
2     const int size_a = n / 2;
3     const int size_b = n - size_a;
4     #pragma omp task final(size_a < 10000) // stop creating tasks if size_a < 10000
5     merge_sort(arr, buffer, size_a); // parallel recursive call to merge sort
6     merge_sort(arr + size_a, buffer + size_a, size_b); // recursive call to merge sort
7     #pragma omp taskwait // wait until both subarrays are sorted
8     merge(arr, arr + size_a, buffer, size_a, size_b, n); // merge from arr into buffer
9 }
10
11 void merge_sort(int *arr, int *buffer, int n) {
12     if (n > 1) { // only this function can end a recursion
13         if (n < 64) { // insertion sort for n smaller than 64
14             insertion_sort(arr, n);
15             return;
16         }
17         const int size_a = n / 2;
18         const int size_b = n - size_a;
19         #pragma omp task final(size_a < 10000) // stop creating tasks if size_a < 10000
20         merge_sort_helper(arr, buffer, size_a); // parallel recursive call to helper
21         merge_sort_helper(arr + size_a, buffer + size_a, size_b); // recursive call to helper
22         #pragma omp taskwait // wait until both subarrays are sorted
23         merge(buffer, buffer + size_a, arr, size_a, size_b, n); // merge from buffer into arr
24     }
25 }
26
27 void merge_sort_run(int *arr, int *buffer, int n) { // buffer is O(n) extra memory
28     #pragma omp parallel // user provides buffer when calling merge_sort_run
29     #pragma omp single nowait
30     merge_sort(arr, buffer, n);
31 }
```

if Clause and final Clause

if(*expr*)

If the expression *expr* evaluates to true, the construct is executed. Can be applied to `#omp parallel` or `#omp task` to trigger parallel execution. *// simplified explanation*

if_clause.cpp ➡

```
1 int main() {  
2     volatile int problem_size = 100;  
3     // if the problem size is small, don't execute in parallel  
4     #pragma omp parallel if (problem_size > 1000)  
5         cout << "Hello from thread " << omp_get_thread_num() << endl;  
6 }
```

final(*expr*)

The final clause takes an expression *expr*. If it evaluates to true, no more tasks are generated and the code is executed immediately. This is propagated to all the child tasks (once a task is final, all child tasks are final too). `final` can only be applied to `#omp task`.

Storage Attributes Can Also Be Used for Tasks

Storage attributes in tasks determine how **variables defined outside the `#omp task` region are handled** ("stored") inside the `#omp task` region.

storage_attributes_task.cpp 

```
1  int main() {
2      int x = 1; // x is shared in a task per default -> outside parallel region
3      #pragma omp parallel num_threads(4) // begin parallel region
4      {
5          int y = 7; // y is firstprivate in a task per default -> inside parallel region
6          #pragma omp single // the code in a single is executed only once
7          {
8              cout << "x: ";
9              for (int i = 0; i < 8; ++i) { // create 8 tasks
10             #pragma omp task firstprivate(x) shared(y) // overwrite storage attributes
11             {
12                 cout << ++x; // prints always 2 because x is firstprivate in each task
13                 ++y; // this y comes from the thread that creates the single
14                 // but y is updated from multiple threads -> possible race condition
15             }
16         }
17     } // implicit barrier of single
18     #pragma omp critical
19     cout << "\ny: " << y; // prints 7 7 7 and (very likely) one 15
20 } // end parallel region
21 }
```

Nested Parallelism

nested_parallelism.cpp ➡

```
1  int main() {
2      omp_set_nested(1); // enable nested parallelism
3      int counter = 0; // counter for total threads executed
4      #pragma omp parallel num_threads(4) // start outer parallel region
5      {
6          // each thread creates two new threads
7          #pragma omp parallel num_threads(2) // start inner parallel region
8          {
9              #pragma omp critical
10             { // thread id in inner parallel region has nothing
11                 // to do with the thread id of the outer parallel region
12                 cout << omp_get_thread_num() << endl; // prints 0 or 1
13                 ++counter;
14             }
15         } // end inner parallel region
16     } // end outer parallel region
17     cout << "counter: " << counter << endl; // prints 8 since 4 * 2 = 8
18 }
```

libstdc++ Parallel Mode

Converts automatically the **standard** (sequential) `std::...` **algorithms** to the appropriate **parallel equivalents**.

[Click here for more information](#)

Compiler flags:

```
-fopenmp -march=native -D_GLIBCXX_PARALLEL
```

Parallel Prefix Sum Computation




FIGURE 1.9

Parallel prefix computation using four processors where each processor is assigned five elements of the input array. **Step 1:** Local summation within each processor; **Step 2:** Prefix sum computation using only the rightmost value of each local array; **Step 3:** Addition of the value computed in Step 2 from the left neighbor to each local array element.

Parallel Programming Concepts and Practice, p. 14

Idea Behind Multithreaded Merging




Figure 27.6 The idea behind the multithreaded merging of two sorted subarrays $T[p_1 \dots r_1]$ and $T[p_2 \dots r_2]$ into the subarray $A[p_3 \dots r_3]$. Letting $x = T[q_1]$ be the median of $T[p_1 \dots r_1]$ and q_2 be the place in $T[p_2 \dots r_2]$ such that x would fall between $T[q_2 - 1]$ and $T[q_2]$, every element in subarrays $T[p_1 \dots q_1 - 1]$ and $T[p_2 \dots q_2 - 1]$ (lightly shaded) is less than or equal to x , and every element in the subarrays $T[q_1 + 1 \dots r_1]$ and $T[q_2 + 1 \dots r_2]$ (heavily shaded) is at least x . To merge, we compute the index q_3 where x belongs in $A[p_3 \dots r_3]$, copy x into $A[q_3]$, and then recursively merge $T[p_1 \dots q_1 - 1]$ with $T[p_2 \dots q_2 - 1]$ into $A[p_3 \dots q_3 - 1]$ and $T[q_1 + 1 \dots r_1]$ with $T[q_2 \dots r_2]$ into $A[q_3 + 1 \dots r_3]$.

Your Paper

- ▶ 3 – 4 pages in ACM sigconf format
<https://www.acm.org/publications/proceedings-template>
- ▶ Example Paper
<http://ceur-ws.org/Vol-2726/paper2.pdf>
- ▶ Basic L^AT_EX template
template_paper.tar.gz ➡
(use BibTeX **not** Biber to compile bibliography)

Coding Warmup

1. Try to **speed up** the **naive merge sort** implementation on page 7.
2. On page 15 you have become acquainted with **multithreaded merging**. On page 18 is an actual **singlethreaded implementation** of the algorithm. **Parallelize** the singlethreaded implementation with tasks.

Parallelise the Merge Function

merge_single_threaded.cpp 

```
1      // TODO: parallelize the first recursive call with a task
2      merge_dac(t1, t2, p1, q1 - 1, p2, q2 - 1, a, p3, __comp);
3      merge_dac(t1, t2, q1 + 1, r1, q2, r2, a, q3 + 1, __comp);
4  }
5  }
6
7  template <class _Type, class _Compare>
8  void parallel_merge(_Type *arr1, _Type *arr2, _Type *out,
9      int64_t size1, int64_t size2, int n_threads, _Compare __comp) {
10     // TODO: parallelize merge_dac, use n_threads as the amount of threads
11     // for the parallel region
12     merge_dac(arr1, arr2, 0, size1 - 1, 0, size2 - 1, out, 0, __comp);
13 }
```

Possible Solution for Parallelising the Merge

merge_multi_threaded.cpp ➡

```
1  #pragma omp task // make the first recursive call a task
2      merge_dac(t1, t2, p1, q1 - 1, p2, q2 - 1, a, p3, __comp);
3      merge_dac(t1, t2, q1 + 1, r1, q2, r2, a, q3 + 1, __comp);
4  }
5  }
6
7  template <class _Type, class _Compare>
8  void parallel_merge(_Type *arr1, _Type *arr2, _Type *out,
9      int64_t size1, int64_t size2, int n_threads, _Compare __comp) {
10     #pragma omp parallel num_threads(n_threads)
11     #pragma omp single nowait
12         merge_dac(arr1, arr2, 0, size1 - 1, 0, size2 - 1, out, 0, __comp);
13 }
```

Exam Assignments

- ▶ Explain how **divide and conquer** algorithms can be parallelized with **tasks** in OpenMP.
- ▶ Describe some **ways** to **speed up merge sort**.
- ▶ What is the idea behind **multithreaded merging**?
- ▶ Do the **coding warmup** on **slide 17**.
- ▶ Read *What every systems programmer should know about concurrency*.

<https://assets.bitbashing.io/papers/concurrency-primer.pdf>

Discuss **two things** you find particularly interesting.